

KNN Find k examples $\{\mathbf{x}^{(i)}, t^{(i)}\}$ closest to the test instance \mathbf{x} and then output majority $\arg \max_{t^z} \sum_{r=1}^k \delta(t^{(z)}, t^{(r)})$. Define $\delta(a, b) = 1$ if $a = b$, 0 otherwise. **Choice of k :** Rule is $k < \sqrt{n}$, small k may overfit, while large may underfit. **Curse of Dim:** In high dimensions, “most” points are approximately the same distance. **Computation Cost:** 0 (minimal) at training/ no learning involved. Query time find N distances in D dimension $\mathcal{O}(ND)$ and $\mathcal{O}(N \log N)$ sorting time.

Entropy $H(X) = -\mathbb{E}_{X \sim p} [\log_2 p(X)] = -\sum_{x \in X} p(x) \log_2 p(x)$ **Multi-class:** $H(X, Y) = -\sum_{x \in X} \sum_{y \in Y} p(x, y) \log_2 p(x, y)$ **Properties:** H is non-negative, $H(Y|X) \leq H(Y)$, $X \perp Y \implies H(Y|X) = H(Y)$, $H(Y|Y) = 0$, and $H(X, Y) = H(X|Y) + H(Y) = H(Y|X) + H(X)$

Expected Conditional Entropy $H(Y|X) = \mathbb{E}_{X \sim p(x)} [H(Y|X)] = \sum_{x \in X} p(x) H(Y|X = x) = -\sum_{x \in X} \sum_{y \in Y} p(x, y) \log_2 p(y|x) = -\mathbb{E}_{(X, Y) \sim p(x, y)} [\log_2 p(Y|X)]$ **Information Gain** $IG(Y|X) = H(Y) - H(Y|X)$

Bias Variance Decomposition Using the square error loss $L(y, t) = \frac{1}{2}(y - t)^2$, **Bias** ($\uparrow \implies$ **underfitting**): How close is our classifier to true target. **Variance** ($\uparrow \implies$ **overfitting**): How widely dispersed are our predictions as we generate new datasets

$$\begin{aligned} \mathbb{E}_{\mathbf{x}, \mathcal{D}} \left[(h_{\mathcal{D}}(\mathbf{x}) - t)^2 \right] &= \mathbb{E}_{\mathbf{x}, \mathcal{D}} \left[(h_{\mathcal{D}}(\mathbf{x}) - \mathbb{E}_{\mathcal{D}} [h_{\mathcal{D}}(\mathbf{x})] + \mathbb{E}_{\mathcal{D}} [h_{\mathcal{D}}(\mathbf{x})] - t)^2 \right] \\ &= \mathbb{E}_{\mathbf{x}, \mathcal{D}} \left[(h_{\mathcal{D}}(\mathbf{x}) - \mathbb{E}_{\mathcal{D}} [h_{\mathcal{D}}(\mathbf{x})])^2 + (\mathbb{E}_{\mathcal{D}} [h_{\mathcal{D}}(\mathbf{x})] - t)^2 + 2(h_{\mathcal{D}}(\mathbf{x}) - \mathbb{E}_{\mathcal{D}} [h_{\mathcal{D}}(\mathbf{x})]) (\mathbb{E}_{\mathcal{D}} [h_{\mathcal{D}}(\mathbf{x})] - t) \right] \\ &= \underbrace{\mathbb{E}_{\mathbf{x}, \mathcal{D}} \left[(h_{\mathcal{D}}(\mathbf{x}) - \mathbb{E}_{\mathcal{D}} [h_{\mathcal{D}}(\mathbf{x})])^2 \right]}_{\text{variance}} + \underbrace{\mathbb{E}_{\mathbf{x}} \left[(\mathbb{E}_{\mathcal{D}} [h_{\mathcal{D}}(\mathbf{x})] - t)^2 \right]}_{\text{bias}} \end{aligned}$$

Bagging with Generating Distribution Suppose we could sample m independent training sets $\{\mathcal{D}_i\}_{i=1}^m$ from p_{dataset} . Learn $h_i := h_{\mathcal{D}_i}$ and our final predictor is $h = 1/m \sum_{i=1}^m h_i$. **Bias Unchanged:** $\mathbb{E}_{\mathcal{D}_1, \dots, \mathcal{D}_m \sim p_{\text{dataset}}} [h(\mathbf{x})] = \frac{1}{m} \sum_{i=1}^m \mathbb{E}_{\mathcal{D}_i \sim p_{\text{dataset}}} [h_i(\mathbf{x})] = \mathbb{E}_{\mathcal{D} \sim p_{\text{dataset}}} [h_{\mathcal{D}}(\mathbf{x})]$ **Variance Reduced:** $\text{Var}_{\mathcal{D}_1, \dots, \mathcal{D}_m} [h(\mathbf{x})] = \frac{1}{m^2} \sum_{i=1}^m \text{Var} [h_i(\mathbf{x})] = \frac{1}{m} \text{Var} [h_{\mathcal{D}}(\mathbf{x})]$

Bootstrap Aggregation Take a single dataset \mathcal{D} with n samples and generate m new datasets, each by sampling n training examples from \mathcal{D} , with replacement. We then average the predictions. We have the reduction in variance to be $\text{Var} \left(\frac{1}{m} \sum_{i=1}^m h_i(\mathbf{x}) \right) = \frac{1}{m} (1 - \rho) \sigma^2 + \rho \sigma^2$

Random Forest Upon bootstrap aggregation, for each bag we choose a random set of features to make the trees grow on (decorrelates predictions, lower ρ).

Bayes Optimality $\mathbb{E}_{\mathbf{x}, \mathcal{D}, t|\mathbf{x}} \left[(h_{\mathcal{D}}(\mathbf{x}) - t)^2 \right] = \underbrace{\mathbb{E}_{\mathbf{x}} \left[(\mathbb{E}_{\mathcal{D}} [h_{\mathcal{D}}(\mathbf{x})] - y_*(\mathbf{x}))^2 \right]}_{\text{bias}} + \underbrace{\mathbb{E}_{\mathbf{x}, \mathcal{D}} \left[(h_{\mathcal{D}}(\mathbf{x}) - \mathbb{E}_{\mathcal{D}} [h_{\mathcal{D}}(\mathbf{x})])^2 \right]}_{\text{variance}} + \underbrace{\mathbb{E}_{\mathbf{x}} [\text{Var}[t|\mathbf{x}]]}_{\text{Bayes}}$

Feature Mapping Some time we want to fit a polynomial curve, we can do this using a feature map $y = \mathbf{w}^\top \boldsymbol{\psi}(x)$ where $\boldsymbol{\psi}(x) = [1, x, x^2, \dots]^\top$. In general the feature map could be anything.

Ridge Regression $\mathbf{w}_{\lambda}^{\text{Ridge}} = \underset{\mathbf{w}}{\text{argmin}} \mathcal{J}_{\text{reg}}(\mathbf{w}) = \underset{\mathbf{w}}{\text{argmin}} \frac{1}{2} \|\mathbf{X}\mathbf{w} - \mathbf{t}\|_2^2 + \frac{\lambda}{2} \|\mathbf{w}\|_2^2 = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{t}$ When $\lambda = 0$ this is just OLS.

Gradient Descent Consider the some cost function \mathcal{J} and we want to optimize it.

- **GD:** $\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{\partial \mathcal{J}}{\partial \mathbf{w}}$; **GD w/ Reg** $\mathbf{w} \leftarrow \mathbf{w} - \alpha \left(\frac{\partial \mathcal{J}}{\partial \mathbf{w}} + \lambda \frac{\partial \mathcal{R}}{\partial \mathbf{w}} \right) = (1 - \alpha \lambda) \mathbf{w} - \alpha \frac{\partial \mathcal{J}}{\partial \mathbf{w}}$
- **mSGD:** Choose mini batch $\mathcal{M} \subset \{1, \dots, N\}$ and update $\mathbf{w} \leftarrow \mathbf{w} - \frac{\alpha}{|\mathcal{M}|} \sum_{i=1}^{|\mathcal{M}|} \frac{\partial \mathcal{L}^{(i)}}{\partial \mathbf{w}}$ **Reasonable size** would be $|\mathcal{M}| \approx 100$
- **SGD:** Choose i at uniform; $\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{\partial \mathcal{L}^{(i)}}{\partial \mathbf{w}}$; **Pro//Cons:** Progress w/o seeing all data//High Variance & Not efficiently vectorized

Cross Entropy Loss $\mathcal{L}_{CE} = -t \log y - (1 - t) \log(1 - y)$ **Logistic CE** $\mathcal{L}_{LCE}(z, t) = \mathcal{L}_{CE}(\sigma(z), t) = t \log(1 + e^{-z}) + (1 - t) \log(1 + e^z)$

Multiclass Classification

- **Softmax Function** Natural generalization of logistic func: $y_k = \text{softmax}(z_1, \dots, z_K)_k = \frac{e^{z_k}}{\sum_{k'} e^{z_{k'}}}$; inputs z_k are called logits.
- **CE Loss, Vectorized** $\mathcal{L}_{\text{CE}}(\mathbf{y}, \mathbf{t}) = -\sum_{k=1}^K t_k \log y_k = -\mathbf{t}^\top (\log \mathbf{y})$ where the log is applied elementwise.
- **Softmax Regression** $\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$, $\mathbf{y} = \text{softmax}(\mathbf{z})$, and $\mathcal{L}_{CE} = -\mathbf{t}^\top (\log \mathbf{y})$; GD Updates is $\mathbf{w}_k \leftarrow \mathbf{w}_k - \alpha \frac{1}{N} \sum_{i=1}^N \left(y_k^{(i)} - t_k^{(i)} \right) \mathbf{x}^{(i)}$ where \mathbf{w}_k means the k -th row of \mathbf{W}

Activation Functions **Identity** $y = z$ **ReLU** $y = \max(0, z)$ **Soft ReLU** $y = \log(1 + e^z)$ **Thresholding** $y = 1$ if $z > 0$ else 0.

Logistic $y = \frac{1}{1 + e^{-z}}$ **tanh** $y = \frac{e^z - e^{-z}}{e^z + e^{-z}}$

Multilayer Perceptron

- **Modularity of Layers** $\mathbf{h}^{(1)} = f^{(1)}(\mathbf{x}) = \phi(\mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)})$, $\mathbf{h}^{(2)} = f^{(2)}(\mathbf{h}^{(1)}) = \phi(\mathbf{W}^{(2)}\mathbf{h}^{(1)} + \mathbf{b}^{(2)})$, \dots , $\mathbf{y} = f^{(L)}(\mathbf{h}^{(L-1)}) = f^{(L)} \circ \dots \circ f^{(1)}(\mathbf{x})$
- **Choice of Last Layer Activation Func** Regression: $\mathbf{y} = f^{(L)}(\mathbf{h}^{(L-1)}) = (\mathbf{w}^{(L)})^\top \mathbf{h}^{(L-1)} + b^{(L)}$; Binary Classification: $\mathbf{y} = f^{(L)}(\mathbf{h}^{(L-1)}) = \sigma \left((\mathbf{w}^{(L)})^\top \mathbf{h}^{(L-1)} + b^{(L)} \right)$
- **Back Propagation** Suppose \mathcal{L} what I want to optimize, then for some variable \mathbf{w} that we want to optimize w.r.t., $\frac{\partial \mathcal{L}}{\partial \mathbf{w}} =: \bar{\mathbf{w}}$
- **Back Prop Cost** Forward: one add-multiplicity operation per weight; Backward: two add-multiplicity operations per weight \implies the Backward pass is about as expensive as two Forward passes. (cost is linear in # of layers, quadratic in # of units per layer)