

# NEURAL NET AND DEEP LEARNING

© by Xia, Tingfeng

Wednesday 12<sup>th</sup> February, 2020

This work is licensed under a Creative Commons “Attribution-NonCommercial-ShareAlike 4.0 International” license.



## Contents

<b>1</b>	<b>Optimization</b>	<b>3</b>
1.1	Prerequisites . . . . .	3
1.2	Determining Curvature . . . . .	3
1.2.1	Definite-ness . . . . .	4
1.3	Convexity . . . . .	4
1.3.1	Convex Sets . . . . .	4
1.3.2	Convex Functions . . . . .	4
1.3.3	Convexity Characterization with Hessian . . . . .	5
1.3.4	Convexity Composition Properties . . . . .	5
1.3.5	Convexity for Linear Models . . . . .	5
1.3.6	Non-Convexity of Networks . . . . .	5
1.4	Nuisance-ness of Cost Surface . . . . .	6
1.4.1	Saddle Point . . . . .	6
1.4.2	Plateaux . . . . .	7
1.4.3	Ill-Conditioned Curvature . . . . .	7
<b>2</b>	<b>Convolutional Neural Net and Image Classification</b>	<b>8</b>
2.1	Motivations for Convolution Layer . . . . .	8
2.2	The Convolution Operator . . . . .	8
2.2.1	1-D Signal Processing . . . . .	8
2.2.2	2-D Convolution . . . . .	9
2.2.3	Properties of Convolution . . . . .	9
2.3	Canonical Kernels . . . . .	9
2.3.1	Blurring Kernel . . . . .	9
2.3.2	Sharpening Kernel . . . . .	9

2.3.3	Edge Detector Kernel . . . . .	9
2.3.4	Vertical Edge Detector Kernel . . . . .	10
2.4	Convolutional Networks . . . . .	10
2.4.1	Pooling Layer . . . . .	10
2.4.2	Non-linearity in Convolutional Layers . . . . .	10
2.4.3	Equivariance and Invariance . . . . .	10
2.4.4	Channels in Convolution Layers . . . . .	11
2.5	Size of Convolutional Neural Nets . . . . .	11
2.5.1	Rule of thumb on size . . . . .	12
2.6	Supervised Pre-training and Transfer Learning . . . . .	12

# 1 Optimization

## 1.1 Prerequisites

**Clairaut's Theorem** states that second order derivatives are such that

$$\frac{\partial^2 \mathcal{J}}{\partial \theta_i \partial \theta_j} = \frac{\partial^2 \mathcal{J}}{\partial \theta_j \partial \theta_i} \quad (1.1)$$

**The Hessian Matrix** is a symmetric matrix (due to Clairaut's Theorem) defined as

$$\mathbf{H} = \nabla^2 \mathcal{J} = \begin{pmatrix} \frac{\partial^2 \mathcal{J}}{\partial \theta_1^2} & \frac{\partial^2 \mathcal{J}}{\partial \theta_1 \partial \theta_2} & \dots & \frac{\partial^2 \mathcal{J}}{\partial \theta_1 \partial \theta_D} \\ \frac{\partial^2 \mathcal{J}}{\partial \theta_2 \partial \theta_1} & \frac{\partial^2 \mathcal{J}}{\partial \theta_2^2} & \dots & \frac{\partial^2 \mathcal{J}}{\partial \theta_2 \partial \theta_D} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 \mathcal{J}}{\partial \theta_D \partial \theta_1} & \frac{\partial^2 \mathcal{J}}{\partial \theta_D \partial \theta_2} & \dots & \frac{\partial^2 \mathcal{J}}{\partial \theta_D^2} \end{pmatrix} \quad (1.2)$$

**Second-Order Taylor Approximation** Locally, a function can be approximated by its second order Taylor approximation. We have, for  $\boldsymbol{\theta}$  that is sufficiently close to  $\boldsymbol{\theta}_0$

$$\mathcal{J}(\boldsymbol{\theta}) \approx \mathcal{J}(\boldsymbol{\theta}_0) + \underbrace{\nabla \mathcal{J}(\boldsymbol{\theta}_0)^\top (\boldsymbol{\theta} - \boldsymbol{\theta}_0)}_{1st \text{ order}} + \underbrace{\frac{1}{2} (\boldsymbol{\theta} - \boldsymbol{\theta}_0)^\top \mathbf{H}(\boldsymbol{\theta}_0) (\boldsymbol{\theta} - \boldsymbol{\theta}_0)}_{quadratic} \quad (1.3)$$

Notice that if  $\boldsymbol{\theta}$  is a critical point, then the gradient would be zero. In such case, the approximation will be

$$\mathcal{J}(\boldsymbol{\theta}) \approx \mathcal{J}(\boldsymbol{\theta}_0) + \frac{1}{2} (\boldsymbol{\theta} - \boldsymbol{\theta}_0)^\top \mathbf{H}(\boldsymbol{\theta}_0) (\boldsymbol{\theta} - \boldsymbol{\theta}_0) \quad (1.4)$$

**Spectral Decomposition** Since the Hessian matrix is symmetric, we know that it is spectral decompose-able, that is

$$\mathbf{H} = \mathbf{Q} \boldsymbol{\Lambda} \mathbf{Q}^\top = \mathbf{Q} \boldsymbol{\Lambda} \mathbf{Q}^{-1} \quad (1.5)$$

where  $\mathbf{Q}$  is the (orthogonal) matrix where columns are the eigenvectors and  $\boldsymbol{\Lambda}$  is the diagonal matrix with eigenvalues on diagonal.

Notice that orthogonal matrices are *necessarily* invertible, with transpose and inverse equal to each other.

## 1.2 Determining Curvature

Often, we refer to the hessian matrix  $\mathbf{H}$  as the curvature of a function. Suppose that we are moving along a line defined by  $\boldsymbol{\theta} + t\mathbf{v}$  for some vector  $\mathbf{v}$ . According to the second order Taylor Approximation, we have

$$\mathcal{J}(\boldsymbol{\theta} + t\mathbf{v}) \approx \mathcal{J}(\boldsymbol{\theta}) + t \nabla \mathcal{J}(\boldsymbol{\theta})^\top \mathbf{v} + \frac{t^2}{2} \mathbf{v}^\top \mathbf{H}(\boldsymbol{\theta}) \mathbf{v} \quad (1.6)$$

Thus, in the case where  $\mathbf{v}^\top \mathbf{H} \mathbf{v} > 0$ , the cost function curves up, i.e. has **positive curvature** where as in the case  $\mathbf{v}^\top \mathbf{H} \mathbf{v} < 0$ , it has a **negative curvature**.

### 1.2.1 Definite-ness

- A matrix  $\mathbf{A}$  is positive definite if  $\mathbf{v}^\top \mathbf{A} \mathbf{v} > 0, \forall \mathbf{v} \neq \mathbf{0}$ .
- We say it is positive semidefinite if  $\mathbf{v}^\top \mathbf{A} \mathbf{v} \geq 0, \forall \mathbf{v} \neq \mathbf{0}$ .
- Equivalently, matrix is positive definite *if and only if* all its eigenvalues are positive, and is positive semidefinite *iff* all its eigenvalues are non-negative. I will now show the case (one direction) for positive definite-ness.

Proof: Let  $\mathbf{v} \neq \mathbf{0} \in \mathbb{R}^n$ , then we know

$$\mathbf{v}^\top \mathbf{A} \mathbf{v} = \mathbf{v}^\top \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^\top \mathbf{v} = (\dagger) \quad (1.7)$$

Now, if we define  $\mathbf{z} = \mathbf{v}^\top \mathbf{Q}$ , we have

$$(\dagger) = \mathbf{z} \mathbf{\Lambda} \mathbf{z}^\top = \begin{bmatrix} z_1 \\ \vdots \\ z_n \end{bmatrix} \begin{bmatrix} \lambda_1 & & \\ & \ddots & \\ & & \lambda_n \end{bmatrix} \begin{bmatrix} z_1 & \dots & z_n \end{bmatrix} = \sum_i \lambda_i z_i^2 \quad (1.8)$$

Since  $\mathbf{v} \neq \mathbf{0}$ , we know  $\mathbf{z} = \mathbf{Q}^\top \mathbf{v} \neq \mathbf{0}$ , we know that  $(\dagger)$  must be positive and this concludes the proof. ■

- For any critical point  $\boldsymbol{\theta}_*$ , if  $\mathbf{H}(\boldsymbol{\theta}_*)$  exists and is positive definite, then  $\boldsymbol{\theta}_*$  is a local minimum.

In naïve words,  
all directions  
curve upwards.

## 1.3 Convexity

### 1.3.1 Convex Sets

Recall that we say set  $\mathcal{S}$  is convex if any  $\mathbf{x}_0, \mathbf{x}_1 \in \mathcal{S}$ , it holds that

$$(1 - \lambda)\mathbf{x}_0 + \lambda\mathbf{x}_1 \in \mathcal{S} \quad \text{for } 0 \leq \lambda \leq 1 \quad (1.9)$$

### 1.3.2 Convex Functions

We say that a function  $f$  is convex if for any  $\mathbf{x}_0, \mathbf{x}_1$  it holds that

$$f((1 - \lambda)\mathbf{x}_0 + \lambda\mathbf{x}_1) \leq (1 - \lambda)f(\mathbf{x}_0) + \lambda f(\mathbf{x}_1) \quad \text{for } 0 \leq \lambda \leq 1 \quad (1.10)$$

### 1.3.3 Convexity Characterization with Hessian

If  $\mathcal{J}$  is twice differentiable, then we can use an equivalent characterization in terms of the Hessian matrix  $\mathbf{H}$

- A twice diff-able function is convex *iff* its Hessian is positive semidefinite everywhere the function is defined.
- In the case of univariate function, where the Hessian would be essentially 1 by 1, we say it is convex *iff* its second derivative is non-negative everywhere.

Possibly add proof for convexity of squared error, logistic CE, and softmax CE.

### 1.3.4 Convexity Composition Properties


- The composition of two convex functions is convex.
- If  $f$  is a non-decreasing univariate function and  $g$  is a convex function, then  $f \circ g$  is convex.

### 1.3.5 Convexity for Linear Models

For a linear model, have the hypothesis  $z = \mathbf{w}^\top \mathbf{x} + b$  is a linear function of  $\mathbf{w}$  and  $b$ . Now, if the loss is convex as a function of  $z$ , then it is convex as a function of  $\mathbf{w}$  and  $b$ . Thus, linear regression, logistic regression, and softmax regression are clearly convex. (by directly applying the rules in the Convexity Composition Properties section above.)

### 1.3.6 Non-Convexity of Networks

Clearly, we want the cost function of interest to be convex, since if that is the case then there will be no *spurious local minima*, meaning that any local minimum will also be a global minimum. In such case, learning is very convenient since gradient based algorithms can just go downhill all the way. Unfortunately, training a network with hidden units cannot be convex due to permutation symmetries – we can re-permute the hidden units in a way that preserves the function computed by the network. An example from slides[1] gives a good intuition of this matter.



## Multilayer Neural Net Non-Convex Argument

- By definition, if a (cost) function  $\mathcal{J}$  is convex, then for any set of points  $\theta_1, \dots, \theta_N$  in its domain, we know

$$\mathcal{J}(\lambda_1 \theta_1 + \dots + \lambda_N \theta_N) \leq \lambda_1 \mathcal{J}(\theta_1) + \dots + \lambda_N \mathcal{J}(\theta_N) \quad \text{for } \lambda_i \geq 0, \sum_i \lambda_i = 1 \quad (1.11)$$

- Due to permutation symmetry, there are  $K!$  (where  $K$  is the number of hidden units in a given hidden layer) permutations of the hidden units in a given layer which all compute the same function.
- Suppose that we average the parameters for all  $K!$  permutations, then we get a degenerate network where all the hidden units are identical (i.e. the network learnt nothing using these identical weights).
- Under the assumption that the cost function was convex, this solution would have to be better than the original one (i.e. the one before we average out all the weights in hidden units) and this leads to a contradiction.<sup>1</sup> —✗— Thus training a multilayer neural nets is non-convex. ■

## 1.4 Nuisance-ness of Cost Surface

In this sub-section, we will discuss how non-nice

### 1.4.1 Saddle Point

A saddle point is a point if

- $\nabla \mathcal{J}(\theta) = \mathbf{0}$ , and
- $\mathbf{H}(\theta)$  has some positive and some negative eigenvalues, i.e. some directions with positive curvature and some with negative curvature.

This could be troublesome since

- If we are exactly on the saddle point, then we are stuck - gradient descent won't get us any further.
- In such case, if we inspect the cost, we would judge that the algorithm has converged. However the solution acquired could be far from optimal.

We can tackle this problem by perturbing the current weight solution slightly in the weight space so that the weight could start moving again. It is important that we don't initialize all of the weights to zero, we should break the symmetry by using small random values.

---

<sup>1</sup>Remark: Generally, it is safe to say that a network that has learnt *absolutely nothing* would have a much higher cost than a network that, at least, converged to a local minimum

### 1.4.2 Plateaux

Plateaux happen more than we think they do, here are some examples that very common

- 0-1 loss
- hard threshold activations
- logistic activations and least square



**Saturated Unit** is when we are in the flat region of the activation function. In the case of logistic non-linearity ( $y = 1/(1 + e^{-x})$ ), when we have very large or very small  $x$ 's the gradient would be very close to zero and thus the weights update will get stuck (or be extremely slow).

**Dead Unit** describes an even more extreme case where the gradient is exactly zero. For example, the negative half of the ReLU activation.

### 1.4.3 Ill-Conditioned Curvature

<sup>2</sup> Below is a typical example[1] of a Ill-conditioned curvature

**Long, narrow ravines:**



As we can see from the gradient updates shown in the level set graph, the solution that we get at each iteration bounces back and forth in the high curvature directions and makes very minimal progress in the low curvature direction.

We can characterize such curvature by noticing that the hessian  $\mathbf{H}$  would have some *large* positive eigenvalues and some eigenvalues close to zero, corresponding to high-curvature and low-curvature directions respectively.

---

<sup>2</sup>This is very common in neural net training.

**Gradient Descent Dynamics** Consider a quadratic objective (which is convex)

$$\mathcal{J}(\boldsymbol{\theta}) = \frac{1}{2} \boldsymbol{\theta}^\top \mathbf{A} \boldsymbol{\theta} \quad \text{where } \mathbf{A} \text{ is +ive semi definite} \quad (1.12)$$

then the gradient descent updates would be

$$\boldsymbol{\theta}_{k+1} \leftarrow \boldsymbol{\theta}_k - \alpha \nabla \mathcal{J}(\boldsymbol{\theta}_k) \quad (1.13)$$

$$= \boldsymbol{\theta}_k - \alpha \mathbf{A} \boldsymbol{\theta}_k \quad (1.14)$$

$$= (\mathbf{I} - \alpha \mathbf{A}) \boldsymbol{\theta}_k \quad (1.15)$$

$$\implies \boldsymbol{\theta}_k = (\mathbf{I} - \alpha \mathbf{A})^k \boldsymbol{\theta}_0 \quad (1.16)$$

## 2 Convolutional Neural Net and Image Classification

### 2.1 Motivations for Convolution Layer

- Images are, usually, large and hence using a lot of fully connected layers would result in an insane amount of parameter to learn and makes calculation intractable.
- In images, there are usually *local* patterns or relationships. For example in the task of semantic segmentation, usually the pixel at bottom left has very little to do with the pixel at top right and using a fully connected layer where every pixel contributes to every output would be wasteful. Thus, we want to have an operation that focus on *local* patterns of the input image.
- Also, the same sorts of features that are useful in analyzing one part of the image will probably be useful for analyzing other parts of the image as well which motivates us in using a “filter” to “slide” across the input.

### 2.2 The Convolution Operator

#### 2.2.1 1-D Signal Processing

Consider two arrays,  $a$  and  $b$ . The result of the convolution will be a new array, where

$$(a * b)_t = \sum_{\tau} a_{\tau} b_{t-\tau} \quad (2.1)$$

Here the summation over  $\tau$  is a lazy notation for saying summing over all combinations that makes sense. <sup>3</sup>

In order to make this indexing notation to work, we have to have the indices start at zero rather than one.

<sup>3</sup>We assume infinite zero padding here, more on this later.



### 2.2.2 2-D Convolution

Consider two two dimensional arrays,  $A$  and  $B$ . The result of the convolution will be such that the slot at  $(A * B)_{ij}$  is calculated as

$$(A * B)_{ij} = \sum_s \sum_t A_{st} B_{i-s, j-t} \quad (2.2)$$

Usually, we call the matrix/tensor that we convolve the original input with a “kernel” or “filter”.

### 2.2.3 Properties of Convolution

- Convolution is *Commutative*, i.e.,

$$a * b = b * a \quad (2.3)$$

- Convolution is *Linear*, i.e.,

$$a * (\lambda_1 b + \lambda_2 c) = \lambda_1 a * b + \lambda_2 a * c \quad (2.4)$$

## 2.3 Canonical Kernels

### 2.3.1 Blurring Kernel

0	1	0
1	4	1
0	1	0

(2.5)

### 2.3.2 Sharpening Kernel

0	-1	0
-1	5	-1
0	-1	0

(2.6)

### 2.3.3 Edge Detector Kernel

0	-1	0
-1	4	-1
0	-1	0

(2.7)

Some of these canonical kernels are not working as expected...

### 2.3.4 Vertical Edge Detector Kernel

$$\begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array} \quad (2.8)$$

## 2.4 Convolutional Networks

In a Convolutional Neural Net, of course there will be convolution layers, however there is another sort of layers that are common, called the pooling layer. Intuitively, pooling layers shrink the dimensions by taking a “local pool”.

### 2.4.1 Pooling Layer

Most commonly, we use the max-pooling operation in the pooling layer, which computes the maximum of the units in a pooling group

$$y_i = \max_{j \text{ in local pooling group}} z_j \quad (2.9)$$

where  $z$  represents the input and  $y$  is the output. Typically, we use a  $2 \times 2$  max pooling unit, which outputs

$$\begin{array}{|c|c|} \hline \alpha & \beta \\ \hline \gamma & \theta \\ \hline \end{array} \longrightarrow \boxed{\max\{\alpha, \beta, \gamma, \theta\}} \quad (2.10)$$

Thanks to pooling layers, deeper layers’ filters will cover a larger region of the input than equal-sized filters in the lower layers. We say that deeper pooling layers have larger receptive fields in terms of the original image.

### 2.4.2 Non-linearity in Convolutional Layers

After convolution operation, it is common to add a non-linear activation function to introduce non-linearity. We are doing this because convolution is a linear operation and stacking convolution layers together without non-linearity is no more powerful than a single linear layer (possibly a fully connected layer). For example, the order of the layers could be

$$\text{Image} \rightarrow \underbrace{\text{Convolution} \rightarrow \text{ReLU activation}}_{\text{Convolution Layer}} \rightarrow \text{Max Pooling} \rightarrow \text{Convolution} \rightarrow \dots \quad (2.11)$$

### 2.4.3 Equivariance and Invariance

TODO: I have no idea...

We want our network’s responses to be robust to translations of the input, which could mean the following two things

- Convolution layers are equivariant<sup>4</sup>: if you translate the inputs, the outputs are translate by the same amount.
- Network’s predictions are invariant: if you translate the inputs, the prediction should not change. Pooling layers provide invariance to small translations.

#### 2.4.4 Channels in Convolution Layers

Each layer is consist of several feature maps, or ***channels***, each of which is an array (and of the same size). In the case where input is an image, then usually we would have 3 channels for coloured input (RGB channels) and 1 channel if it is in greyscale. Each unit is connected to each unit within its receptive field in the previous layer. This includes *all* of the previous layer’s feature maps.

### 2.5 Size of Convolutional Neural Nets


There are several “measures of sizes” that are interested in, namely

- ***Number of units***: measures the activations needed to be stored in memory during training for back propagation.
- ***Number of weights***: since weights need to be stored (and updated at each iteration) in memory.
- ***Number of connections***: measures the computation costs; approximately 3 add-multiply operations per connection. (1 for the forward pass, and 2 for backward pass.)

Below is the number of parameters in a fully connected multilayer perceptron layer along side with those in a convolutional layer. [2]

---

<sup>4</sup>Equivariant means roughly “unchanged in terms of distortion”



	fully connected layer	convolution layer
# output units	$WHI$	$WHI$
# weights	$W^2H^2IJ$	$K^2IJ$
# connections	$W^2H^2IJ$	$WHK^2IJ$

### 2.5.1 Rule of thumb on size

- Most of the units and connections are in the convolution layers, and
- Most of the weights are in the fully connected layers.

## 2.6 Supervised Pre-training and Transfer Learning

- In practice, we don't usually train an image classifier from scratch since it is unlikely that we will have millions of cleanly labeled images for our specific datasets.
- If we want to do a computer vision task, it is common to just "fine-tune" a pre-trained convolutional neural net on ImageNet or OpenImage.
- We will fix most of the weights in the pre-trained network. Only the weights in the last layer will be randomly initialized and learnt on the current dataset/task.

### Fine-tune, how?

- This depends on how many training examples do we have in the new dataset. For example, if we have fewer examples, then we can fix more weights from the pre-trained networks and only train a small subset of all weights. Of course, vice versa.

- This also depends on how similar is the new dataset to the dataset that our pre-trained model was trained on. That is, we need ‘more’ fine-tuning if they have dissimilar datasets.
- **#! Important:** Learning rate for the fine-tuning stage is often much lower than the learning rate used for training from scratch. \_\_\_\_\_

why?

## References

- [1] Jimmy Ba. Csc413/2516 lecture 4: Optimization. 2020 Winter Term.
- [2] Jimmy Ba. Csc421/2516 lecture 5: Convolutional neural network & image classification. 2020 Winter Term.